

DATASCI W261: Machine Learning at Scale

Assignment Week 3

Minhchau Dang (minhchau.dang@berkeley.edu)

Miki Seltzer (miki.seltzer@berkeley.edu)

W261-2, Spring 2016

Submission: 2016-Feb-02 8:50pm

HW3.0

What is a merge sort? Where is it used in Hadoop?

Merge sort is used to combine two pre-sorted lists. It is a very efficient sort, as it only needs to iteratively look for the smallest element in multiple sorted lists. It is utilized in Hadoop during the shuffle, when key-value pairs are shuffled to reducers, then sorted.

How is a combiner function in the context of Hadoop?

A combiner function allows for preaggregation before key-value pairs are sent from the mappers to reducers. It is similar to a reducer, but there are some key differences. One difference is that the combiner may not always be used during the job -- it may be used 0, 1, or many times. Thus, we cannot count on Hadoop actually using a combiner we have included in the job, and we must be careful about matching output types of the mapper and combiner.

Give an example where it can be used and justify why it should be used in the context of this problem.

In the classic word count example, a document is scanned, and each word is paired with the value of 1. A combiner can be used to combine values of 1 with the same key (word) before they are shuffled to reducers. This reduces the amount of data that is shuffled between the mapper and reducer, and increases efficiency.

What is the Hadoop shuffle?

The Hadoop shuffle is the process by which data from mappers is shuffled and sorted while being sent to reducers. The shuffle ensures that keys are grouped together and sorted within the reducer they are sent to.

HW3.1: Use Counters to do EDA (exploratory data analysis and to monitor progress)

```
In [343]: # I am running this locally, so make sure that the Hadoop streaming API
           is in this folder.
           !wget http://central.maven.org/maven2/org/apache/hadoop/hadoop-streamin
           g/2.7.1/hadoop-streaming-2.7.1.jar

           # Create a folder on HDFS for this week's assignment, strip the header l
           ine from Consumer_Complaints.csv
           !echo "${tail -n +2 Consumer_Complaints.csv}" > Consumer_Complaints.csv
           !hdfs dfs -mkdir /user/miki/week03
           !hdfs dfs -put Consumer_Complaints.csv /user/miki/week03
```

mkdir: `/user/miki/week03': File exists

```
In [344]: %%writefile mapper_31.py
           #!/usr/bin/python
           ## mapper.py
           ## Author: Miki Seltzer
           ## Description: mapper code for HW3.1

           import sys
           from csv import reader

           # Our input comes from STDIN (standard input)
           for line in reader(sys.stdin):
               product = line[1]
               if product == "Debt collection": sys.stderr.write("reporter:counter:
           Product,Debt,1\n")
               elif product == "Mortgage": sys.stderr.write("reporter:counter:Produ
           ct,Mortgage,1\n")
               else: sys.stderr.write("reporter:counter:Product,Other,1\n")
               print line
```

Writing mapper_31.py

```
In [345]: %%writefile reducer_31.py
#!/usr/bin/python
## reducer.py
## Author: Miki Seltzer
## Description: reducer code for HW3.1

import sys
from operator import itemgetter
from csv import reader

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    print line
```

Writing reducer_31.py

```
In [346]: # Change permissions on mapper and reducer
!chmod +x mapper_31.py
!chmod +x reducer_31.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_1_output

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_31.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_31.py \
-input /user/miki/week03/Consumer_Complaints.csv \
-output /user/miki/week03/hw3_1_output

Deleted /user/miki/week03/hw3_1_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob7547497993532515688.jar tmpDir=null
```

Screenshot



HW 3.2 Analyze the performance of your Mappers, Combiners and Reducers using Counters

```
In [18]: %%writefile HW3_2_input.txt
foo foo quux labs foo bar quux
```

Overwriting HW3_2_input.txt

```
In [19]: !hdfs dfs -put HW3_2_input.txt /user/miki/week03
```

```
In [347]: %%writefile mapper_32a.py
#!/usr/bin/python
## mapper.py
## Author: Miki Seltzer
## Description: mapper code for HW3.2

import sys
from csv import reader

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Mapper,1\n")

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    line = line.split()
    for word in line:
        print '%s\t%s' % (word, 1)
```

Writing mapper_32a.py

```
In [348]: %%writefile reducer_32a.py
#!/usr/bin/python
## reducer.py
## Author: Miki Seltzer
## Description: reducer code for HW3.2

import sys
from operator import itemgetter

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Reducer,1\n")

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    print line
```

Writing reducer_32a.py

```
In [349]: # Change permissions on mapper and reducer
!chmod +x mapper_32a.py
!chmod +x reducer_32a.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_2a_output

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.map.tasks=1 \
-D mapred.reduce.tasks=4 \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_32a.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_32a.py \
-input /user/miki/week03/HW3_2_input.txt \
-output /user/miki/week03/hw3_2a_output

Deleted /user/miki/week03/hw3_2a_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob7785257004505429063.jar tmpDir=null
```

What is the value of your user defined Mapper Counter, and Reducer Counter after completing this word count job? The answer should be 1 and 4 respectively. Please explain.

I had to specify the number of map tasks and reduce tasks to get 1 and 4, since the defaults produced counters of 2 and 1 respectively.

The counters were incremented each time the mapper and reducer scripts were executed.

HW3.2b: Perform a word count analysis of the Issue column of the Consumer Complaints Dataset

```
In [302]: %%writefile mapper_32b.py
          #!/usr/bin/python
          ## mapper.py
          ## Author: Miki Seltzer
          ## Description: mapper code for HW3.2b

          import sys
          from csv import reader
          import string

          # Increment mapper counter
          sys.stderr.write("reporter:counter:Custom_Counter,Mapper,1\n")

          # Initialize variables
          total = 0

          # Our input comes from STDIN (standard input)
          for line in reader(sys.stdin):
              # Format our line
              issue = line[3].lower()
              issue = issue.replace(',', ' ').replace('/', ' ')

              for word in issue.split():
                  if len(word) > 0:
                      print '%s\t%s' % (word, 1)
                      total += 1

          # Print total words
          print '%s\t%s' % ('*total', total)
```

Writing mapper_32b.py

```
In [303]: %%writefile reducer_32b.py
#!/usr/bin/python
## reducer.py
## Author: Miki Seltzer
## Description: reducer code for HW3.2b

import sys
from operator import itemgetter

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Reducer,1\n")

# Initialize variables
prev_word = None
prev_count = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:

    # Split line
    word, count = line.split('\t')

    # Convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # Count wasn't an int, so move on
        continue

    if prev_word == word:
        # We haven't moved to a new word
        prev_count += count

    else:
        if prev_word:
            print '%s\t%s' % (prev_word, prev_count)

        prev_count = count
        prev_word = word

# Output the last line
if prev_word == word:
    print '%s\t%s' % (prev_word, prev_count)
```

Writing reducer_32b.py

```
In [350]: # Change permissions on mapper and reducer
!chmod +x mapper_32b.py
!chmod +x reducer_32b.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_2b_output

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.reduce.tasks=4 \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_32b.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_32b.py \
-input /user/miki/week03/Consumer_Complaints.csv \
-output /user/miki/week03/hw3_2b_output
```

```
Deleted /user/miki/week03/hw3_2b_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob1465294680810041979.jar tmpDir=null
```

What is the value of your user defined Mapper Counter, and Reducer Counter after completing your word count job?

After completing this job, the counters show the following values:

- Mapper: 2
- Reducer: 4 (this is explicitly set when running the job)

HW3.2c: Perform a word count analysis of the Issue column of the Consumer Complaints Dataset (ADD: standalone combiner)

We can reuse the reducer in this case, and rename it combiner. We update the line to increment the combiner counter.


```
In [305]: %%writefile combiner_32c.py
#!/usr/bin/python
## combiner.py
## Author: Miki Seltzer
## Description: combiner code for HW3.2c

import sys
from operator import itemgetter

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Combiner,1\n")

prev_word = None
prev_count = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    word, count = line.split('\t')

    # Convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # Count wasn't an int, so move on
        continue

    # Check if we've moved to a new word
    if prev_word == word:
        prev_count += count
    else:
        if prev_word:
            # We are at a new word, need to print previous word sum
            print '%s\t%s' % (prev_word, prev_count)
        prev_count = count
        prev_word = word

# Output the last line
if prev_word == word:
    print '%s\t%s' % (prev_word, prev_count)
```

Writing combiner_32c.py

```
In [351]: # Change permissions on mapper and reducer
!chmod +x mapper_32b.py
!chmod +x combiner_32c.py
!chmod +x reducer_32b.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_2c_output

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.reduce.tasks=4 \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_32b.py \
-combiner /home/cloudera/Documents/W261-Fall2016/Week03/combiner_32c.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_32b.py \
-input /user/miki/week03/Consumer_Complaints.csv \
-output /user/miki/week03/hw3_2c_output

Deleted /user/miki/week03/hw3_2c_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob2612824031763828089.jar tmpDir=null
```

What is the value of your user defined Mapper Counter, Combiner Counter and Reducer Counter after completing your word count job?

After completing this job, the counters show the following values:

- Mapper: 2
- Combiner: 8
- Reducer: 4 (this is explicitly set when running the job)

HW3.2d: Using a single reducer, present frequency and relative frequency of top 50 and bottom 10 terms

For this section, we only need an identity mapper and an identity reducer.

```
In [309]: %%writefile mapper_32d.py
#!/usr/bin/python
## mapper.py
## Author: Miki Seltzer
## Description: mapper code for HW3.2d

import sys
from csv import reader
import string

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Mapper,1\n")

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    word, count = line.replace('\n','').split('\t')
    print '%s\t%s' % (count, word)
```

Overwriting mapper_32d.py

```
In [310]: %%writefile reducer_32d.py
#!/usr/bin/python
## reducer.py
## Author: Miki Seltzer
## Description: reducer code for HW3.2d

import sys
from operator import itemgetter

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Reducer,1\n")

# Initialize variables
total = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    fields = line.replace('\n','').split('\t')
    count = fields[0]
    word = fields[1]

    try:
        count = int(count)
    except ValueError:
        continue

    # The first word should be *total, save this as total
    if word == '*total': total = float(count)
    else: print '%s\t%s\t%s' % (word, count, count/total)
```

Overwriting reducer_32d.py

```
In [352]: # Change permissions on mapper and reducer
!chmod +x mapper_32d.py
!chmod +x reducer_32d.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_2d_output

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFiel
eldBasedComparator \
-D mapred.text.key.partitioner.options=-k1,1 \
-D stream.num.map.output.key.fields=2 \
-D mapred.text.key.comparator.options='-k1,1nr -k2,2n' \
-D mapred.reduce.tasks=1 \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_32d.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_32d.py \
-input /user/miki/week03/hw3_2b_output/part* \
-output /user/miki/week03/hw3_2d_output

Deleted /user/miki/week03/hw3_2d_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob6548003399415369612.jar tmpDir=null
```

```
In [353]: !rm hw3_2d_output.txt
!hdfs dfs -copyToLocal /user/miki/week03/hw3_2d_output/part-00000 hw3_2d
_output.txt
```

```

In [354]: # Function to pretty print:
# - the top x and bottom y items
# - unique items
def print_results(file, x=50, y=10):
    words = []
    special_words = []

    with open(file, 'r') as myfile:
        for line in myfile:
            fields = line.replace('\n', '').split('\t')
            if fields[0][0] != '*': words.append(fields)
            else: special_words.append(fields)

    print '          {:16s}{: >8s}{: >15s}'.format('word', 'count', 'relative
freq')
    print '-----'
    for i in range(x):
        print '[:3d]  {:16s}{:8,d}{:15.2%}'.format(i+1,
                                                    words[i][0],
                                                    int(words[i][1]),
                                                    float(words[i][2]))

    print '...'
    for i in range(y):
        j = len(words) - 10 + i
        print '[:3d]  {:16s}{:8,d}{:15.2%}'.format(j+1,
                                                    words[j][0],
                                                    int(words[j][1]),
                                                    float(words[j][2]))

    print '\n-----'
    print '          {:16s}{: >8,d}'.format('Unique words', len(words))
    for item in special_words:
        name = item[0][1:].replace('_', ' ')
        print '          {:16s}{: >8,d}'.format(name, int(item[1]))

print_results('hw3_2d_output.txt')

```


	word	count	relative freq
[1]	loan	119,630	8.87%
[2]	collection	72,394	5.37%
[3]	foreclosure	70,487	5.23%
[4]	modification	70,487	5.23%
[5]	account	57,448	4.26%
[6]	credit	55,251	4.10%
[7]	or	40,508	3.00%
[8]	payments	39,993	2.97%
[9]	escrow	36,767	2.73%
[10]	servicing	36,767	2.73%
[11]	report	34,903	2.59%
[12]	incorrect	29,133	2.16%
[13]	information	29,069	2.16%
[14]	on	29,069	2.16%
[15]	debt	27,874	2.07%
[16]	closing	19,000	1.41%
[17]	not	18,477	1.37%
[18]	attempts	17,972	1.33%
[19]	cont'd	17,972	1.33%
[20]	collect	17,972	1.33%
[21]	owed	17,972	1.33%
[22]	and	16,448	1.22%
[23]	management	16,205	1.20%
[24]	opening	16,205	1.20%
[25]	of	13,983	1.04%
[26]	my	10,731	0.80%
[27]	deposits	10,555	0.78%
[28]	withdrawals	10,555	0.78%
[29]	problems	9,484	0.70%
[30]	application	8,868	0.66%
[31]	communication	8,671	0.64%
[32]	tactics	8,671	0.64%
[33]	broker	8,625	0.64%
[34]	mortgage	8,625	0.64%
[35]	originator	8,625	0.64%
[36]	to	8,401	0.62%
[37]	unable	8,178	0.61%
[38]	billing	8,158	0.61%
[39]	other	7,886	0.58%
[40]	disclosure	7,655	0.57%
[41]	verification	7,655	0.57%
[42]	disputes	6,938	0.51%
[43]	reporting	6,559	0.49%
[44]	lease	6,337	0.47%
[45]	the	6,248	0.46%
[46]	by	5,663	0.42%
[47]	being	5,663	0.42%
[48]	caused	5,663	0.42%
[49]	funds	5,663	0.42%
[50]	low	5,663	0.42%
...			

[165]	apply	118	0.01%
[166]	amount	98	0.01%
[167]	credited	92	0.01%
[168]	payment	92	0.01%
[169]	checks	75	0.01%
[170]	convenience	75	0.01%
[171]	amt	71	0.01%
[172]	day	71	0.01%
[173]	disclosures	64	0.00%
[174]	missing	64	0.00%

Unique words	174
--------------	-----

HW3.3. Shopping Cart Analysis Exploratory Data Analysis

We can reuse the reducer from HW3.2b, but there are small changes that need to be made to the mapper:

- We do not have to format the products to lower case, assume there is no punctuation stripping needed
- Keep track of the largest basket size as we loop through baskets

In [282]: `!hdfs dfs -put ProductPurchaseData.txt /user/miki/week03`

```
In [325]: %%writefile mapper_33a.py
#!/usr/bin/python
## mapper.py
## Author: Miki Seltzer
## Description: mapper code for HW3.3a

import sys

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Mapper,1\n")

# Initialize variables
total = 0
basket_size = 0
largest_basket_size = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    # Split our line into products
    for product in line.replace('\n','').split():
        print '%s\t%s' % (product, 1)
        basket_size += 1
        total += 1
    if basket_size > largest_basket_size:
        largest_basket_size = basket_size

    basket_size = 0

# Print total words
print '%s\t%s' % (*total, total)
print '%s\t%s' % (*largest_basket, largest_basket_size)
```

Overwriting mapper_33a.py

```
In [326]: # Change permissions on mapper and reducer
!chmod +x mapper_33a.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_3a_output

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_33a.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_32b.py \
-input /user/miki/week03/ProductPurchaseData.txt \
-output /user/miki/week03/hw3_3a_output
```

```
Deleted /user/miki/week03/hw3_3a_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob4662501162287957069.jar tmpDir=null
```

Using a single reducer: Report your findings such as number of unique products; largest basket; report the top 50 most frequently purchased items, their frequency, and their relative frequency (break ties by sorting the products alphabetical order) etc. using Hadoop Map-Reduce.

We can use the mapper and reducer from HW3.2d to get the sorted frequencies and relative frequencies of the products.

```
In [327]: # If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_3b_output

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFi
eldBasedComparator \
-D mapred.text.key.partitionner.options=-k1,1 \
-D stream.num.map.output.key.fields=2 \
-D mapred.text.key.comparator.options='-k1,1nr -k2,2n' \
-D mapred.reduce.tasks=1 \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_32d.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_32d.py \
-input /user/miki/week03/hw3_3a_output/part* \
-output /user/miki/week03/hw3_3b_output
```

```
Deleted /user/miki/week03/hw3_3b_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob7461671848105384853.jar tmpDir=null
```

```
In [330]: !rm hw3_3b_output.txt  
          !hdfs dfs -copyToLocal /user/miki/week03/hw3_3b_output/part-00000 hw3_3b  
          _output.txt
```

```
In [341]: print_results('hw3_3b_output.txt', 50, 0)
```


	word	count	relative freq

[1]	DAI62779	6,667	1.75%
[2]	FR040251	3,881	1.02%
[3]	ELE17451	3,875	1.02%
[4]	GRO73461	3,602	0.95%
[5]	SNA80324	3,044	0.80%
[6]	ELE32164	2,851	0.75%
[7]	DAI75645	2,736	0.72%
[8]	SNA45677	2,455	0.64%
[9]	FR031317	2,330	0.61%
[10]	DAI85309	2,293	0.60%
[11]	ELE26917	2,292	0.60%
[12]	FR080039	2,233	0.59%
[13]	GRO21487	2,115	0.56%
[14]	SNA99873	2,083	0.55%
[15]	GRO59710	2,004	0.53%
[16]	GRO71621	1,920	0.50%
[17]	FR085978	1,918	0.50%
[18]	GRO30386	1,840	0.48%
[19]	ELE74009	1,816	0.48%
[20]	GRO56726	1,784	0.47%
[21]	DAI63921	1,773	0.47%
[22]	GRO46854	1,756	0.46%
[23]	ELE66600	1,713	0.45%
[24]	DAI83733	1,712	0.45%
[25]	FR032293	1,702	0.45%
[26]	ELE66810	1,697	0.45%
[27]	SNA55762	1,646	0.43%
[28]	DAI22177	1,627	0.43%
[29]	FR078087	1,531	0.40%
[30]	ELE99737	1,516	0.40%
[31]	ELE34057	1,489	0.39%
[32]	GRO94758	1,489	0.39%
[33]	FR035904	1,436	0.38%
[34]	FR053271	1,420	0.37%
[35]	SNA93860	1,407	0.37%
[36]	SNA90094	1,390	0.36%
[37]	GRO38814	1,352	0.36%
[38]	ELE56788	1,345	0.35%
[39]	GRO61133	1,321	0.35%
[40]	DAI88807	1,316	0.35%
[41]	ELE74482	1,316	0.35%
[42]	ELE59935	1,311	0.34%
[43]	SNA96271	1,295	0.34%
[44]	DAI43223	1,290	0.34%
[45]	ELE91337	1,289	0.34%
[46]	GRO15017	1,275	0.33%
[47]	DAI31081	1,261	0.33%
[48]	GRO81087	1,220	0.32%
[49]	DAI22896	1,219	0.32%
[50]	GRO85051	1,214	0.32%

...

Unique words	12,592
largest basket	74

HW3.4: Pairs

Suppose we want to recommend new products to the customer based on the products they have already browsed on the online website. Write a map-reduce program to find products which are frequently browsed together. Fix the support count (cooccurrence count) to $s = 100$ (i.e. product pairs need to occur together at least 100 times to be considered frequent) and find pairs of items (sometimes referred to itemsets of size 2 in association rule mining) that have a support count of 100 or more.

```
In [360]: %%writefile mapper_34a.py
#!/usr/bin/python
## mapper.py
## Author: Miki Seltzer
## Description: mapper code for HW3.4a

import sys
import itertools

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Mapper,1\n")

# Initialize variables
total = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    # Split our line into products
    products = line.replace('\n','').split()

    # Get all combinations of products:
    # - Use a set to remove duplicate products
    # - Combinations finds tuples of length 2 with no repeats
    pairs = list(itertools.combinations(set(products), 2))

    # For each pair, sort the pair alphabetically, then emit
    for pair in pairs:
        sorted_pair = sorted(pair)
        print '%s\t%s\t%s' % (sorted_pair[0], sorted_pair[1], 1)

    # Increment total number of baskets
    total += 1

# Print total words
print '%s\t%s\t%s' % (*total, '', total)
```

Overwriting mapper_34a.py

```

In [365]: %%writefile reducer_34a.py
#!/usr/bin/python
## reducer.py
## Author: Miki Seltzer
## Description: reducer code for HW3.4a

import sys
from operator import itemgetter

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Reducer,1\n")

# Initialize variables
prev_pair = []
prev_count = 0
total = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    # Define our key and value
    fields = line.replace('\n', '').split('\t')
    pair = [fields[0], fields[1]]
    count = fields[2]

    # Convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # Count wasn't an int, so move on
        continue

    # Check if we've moved to a new word
    if prev_pair == pair:
        prev_count += count
    else:
        if len(prev_pair) > 0:
            # We are at a new pair, need to print previous pair sum
            print '%s\t%s\t%s' % (prev_pair[0], prev_pair[1], prev_count)

        prev_count = count
        prev_pair = pair

# Output the last line
if prev_pair == pair:
    print '%s\t%s\t%s' % (prev_pair[0], prev_pair[1], prev_count)

```

Overwriting reducer_34a.py

```
In [21]: # Change permissions on mapper and reducer
!chmod +x mapper_34a.py
!chmod +x reducer_34a.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_4a_output

# Run job
!time hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFi
eldBasedComparator \
-D mapred.text.key.partitionner.options=-k1,1 \
-D stream.num.map.output.key.fields=2 \
-D mapred.text.key.comparator.options='-k1,1 -k2,2' \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_34a.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_34a.py \
-input /user/miki/week03/ProductPurchaseData.txt \
-output /user/miki/week03/hw3_4a_output

Deleted /user/miki/week03/hw3_4a_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob5023424855138306517.jar tmpDir=null

real    0m50.284s
user    0m5.474s
sys     0m0.365s
```

Now we have each pair and the number of times that the pair co-occurs. We need to run another job to calculate the relative frequency and sort the resulting pairs

```
In [379]: %%writefile mapper_34b.py
#!/usr/bin/python
## mapper.py
## Author: Miki Seltzer
## Description: mapper code for HW3.4b

import sys

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Mapper,1\n")

# Initialize variables
total = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    fields = line.replace('\n','').split('\t')
    if int(fields[2]) >= 100:
        print '%s\t%s\t%s' % (fields[2], fields[0], fields[1])
```

Overwriting mapper_34b.py

```
In [381]: %%writefile reducer_34b.py
#!/usr/bin/python
## reducer.py
## Author: Miki Seltzer
## Description: reducer code for HW3.4b

import sys
from operator import itemgetter

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Reducer,1\n")

# Initialize variables
total = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    fields = line.replace('\n','').split('\t')
    count = fields[0]
    item1 = fields[1]
    item2 = fields[2]

    try:
        count = int(count)
    except ValueError:
        continue

    # The first word should be *total, save this as total
    if item1 == '*total': total = float(count)
    else: print '%s\t%s\t%s\t%s' % (item1, item2, count, count/total)
```

Writing reducer_34b.py

```
In [383]: # Change permissions on mapper and reducer
!chmod +x mapper_34b.py
!chmod +x reducer_34b.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_4b_output

# Run job
!time hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFi
eldBasedComparator \
-D mapred.text.key.partitioner.options=-k1,1 \
-D stream.num.map.output.key.fields=3 \
-D mapred.text.key.comparator.options='-k1,1nr -k2,2 -k3,3' \
-D mapred.reduce.tasks=1 \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_34b.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_34b.py \
-input /user/miki/week03/hw3_4a_output/part* \
-output /user/miki/week03/hw3_4b_output
```

```
Deleted /user/miki/week03/hw3_4b_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob2299545763059091234.jar tmpDir=null
```

```
real    0m26.758s
user    0m4.982s
sys     0m0.345s
```

```
In [386]: !hdfs dfs -copyToLocal /user/miki/week03/hw3_4b_output/part-00000 hw3_4b
_output.txt
```

```

In [31]: # Function to pretty print:
# - the top x and bottom y items
# - unique items
def print_results(file, x=50, y=10):
    words = []
    special_words = []

    with open(file, 'r') as myfile:
        for line in myfile:
            fields = line.replace('\n', '').split('\t')
            if fields[0][0] != '*': words.append(fields)
            else: special_words.append(fields)

    print '          {:10s}{:10s}{:>8s}{:>15s}'.format('item1', 'item2', 'count', 'relative freq')
    print '-----'
    for i in range(x):
        print '[:3d]  {:10s}{:10s}{:8,d}{:15.2%}'.format(i+1,
                                                         words[i][0],
                                                         words[i][1],
                                                         int(words[i]
                                                         [2]),
                                                         float(words[i]
                                                         [3]))

```

List the top 50 product pairs


```
In [395]: print_results('hw3_4b_output.txt', 50, 0)
```

	item1	item2	count	relative freq
[1]	DAI62779	ELE17451	1,592	5.12%
[2]	FR040251	SNA80324	1,412	4.54%
[3]	DAI75645	FR040251	1,254	4.03%
[4]	FR040251	GR085051	1,213	3.90%
[5]	DAI62779	GR073461	1,139	3.66%
[6]	DAI75645	SNA80324	1,130	3.63%
[7]	DAI62779	FR040251	1,070	3.44%
[8]	DAI62779	SNA80324	923	2.97%
[9]	DAI62779	DAI85309	918	2.95%
[10]	ELE32164	GR059710	911	2.93%
[11]	DAI62779	DAI75645	882	2.84%
[12]	FR040251	GR073461	882	2.84%
[13]	DAI62779	ELE92920	877	2.82%
[14]	FR040251	FR092469	835	2.68%
[15]	DAI62779	ELE32164	832	2.68%
[16]	DAI75645	GR073461	712	2.29%
[17]	DAI43223	ELE32164	711	2.29%
[18]	DAI62779	GR030386	709	2.28%
[19]	ELE17451	FR040251	697	2.24%
[20]	DAI85309	ELE99737	659	2.12%
[21]	DAI62779	ELE26917	650	2.09%
[22]	GR021487	GR073461	631	2.03%
[23]	DAI62779	SNA45677	604	1.94%
[24]	ELE17451	SNA80324	597	1.92%
[25]	DAI62779	GR071621	595	1.91%
[26]	DAI62779	SNA55762	593	1.91%
[27]	DAI62779	DAI83733	586	1.88%
[28]	ELE17451	GR073461	580	1.86%
[29]	GR073461	SNA80324	562	1.81%
[30]	DAI62779	GR059710	561	1.80%
[31]	DAI62779	FR080039	550	1.77%
[32]	DAI75645	ELE17451	547	1.76%
[33]	DAI62779	SNA93860	537	1.73%
[34]	DAI55148	DAI62779	526	1.69%
[35]	DAI43223	GR059710	512	1.65%
[36]	ELE17451	ELE32164	511	1.64%
[37]	DAI62779	SNA18336	506	1.63%
[38]	ELE32164	GR073461	486	1.56%
[39]	DAI62779	FR078087	482	1.55%
[40]	DAI85309	ELE17451	482	1.55%
[41]	DAI62779	GR094758	479	1.54%
[42]	DAI62779	GR021487	471	1.51%
[43]	GR085051	SNA80324	471	1.51%
[44]	ELE17451	GR030386	468	1.50%
[45]	FR085978	SNA95666	463	1.49%
[46]	DAI62779	FR019221	462	1.49%
[47]	DAI62779	GR046854	461	1.48%
[48]	DAI43223	DAI62779	459	1.48%
[49]	ELE92920	SNA18336	455	1.46%
[50]	DAI88079	FR040251	446	1.43%

Report the compute time for the Pairs job.

The 1st job (counts) reports the following compute times:

real	0m50.284s
user	0m5.474s
sys	0m0.365s

The 2nd job (sorts) reports the following compute times:

real	0m26.758s
user	0m4.982s
sys	0m0.345s

Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers)

Cloudera QuickStart VM: single computer, 2 processors, 2 mappers (default), 1 reducer

How many times is each mapper and reducer called?

- Mapper: 2
- Reducer: 1

HW3.5: Stripes

Repeat 3.4 using the stripes design pattern for finding cooccurring pairs.

```
In [5]: %%writefile mapper_35a.py
#!/usr/bin/python
## mapper.py
## Author: Miki Seltzer
## Description: mapper code for HW3.5a

import sys
import itertools

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Mapper,1\n")

# Initialize variables
total = 0
stripes = {}

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    # Split our line into products
    products = line.replace('\n', '').split()

    # Get all combinations of products:
    # - Use a set to remove duplicate products
    # - Combinations finds tuples of length 2 with no repeats
    items = sorted(list(set(products)))

    for i in range(len(items)-1):
        for j in range(i+1, len(items)):
            stripes[items[j]] = 1
            print '%s\t%s' % (items[i], stripes)
            stripes = {}

    # Increment total number of baskets
    total += 1

# Print total words
print '%s\t%s' % ('*total', {'*total':total})
```

Overwriting mapper_35a.py

```
In [9]: %%writefile reducer_35a.py
#!/usr/bin/python
## reducer.py
## Author: Miki Seltzer
## Description: reducer code for HW3.5a

import sys
from operator import itemgetter

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Reducer,1\n")

# Initialize variables
prev_word = None
prev_stripe = {}
total = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    # Define our key and value
    fields = line.replace('\n','').split('\t')
    word = fields[0]
    stripe = eval(fields[1])

    # Check if we've moved to a new word
    if prev_word == word:
        # We need to move through the dictionary and update counts
        for item in stripe:
            if item in prev_stripe:
                prev_stripe[item] += stripe[item]
            else:
                prev_stripe[item] = stripe[item]
    else:
        if len(prev_stripe) > 0:
            # We are at a new pair, need to print previous pair sum
            print '%s\t%s' % (prev_word, prev_stripe)
            prev_stripe = stripe
            prev_word = word

# Output the last line
if prev_stripe == stripe:
    print '%s\t%s' % (prev_word, prev_stripe)
```

Overwriting reducer_35a.py

```
In [22]: # Change permissions on mapper and reducer
!chmod +x mapper_35a.py
!chmod +x reducer_35a.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_5a_output

# Run job
!time hadoop jar hadoop-streaming-2.7.1.jar \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_35a.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_35a.py \
-input /user/miki/week03/ProductPurchaseData.txt \
-output /user/miki/week03/hw3_5a_output
```

```
Deleted /user/miki/week03/hw3_5a_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob8692232309490935031.jar tmpDir=null
```

```
real    0m44.977s
user    0m5.111s
sys     0m0.337s
```

Now we have a stripe for each product (item1). The stripe contains the second item in the pair (item2), along with the count of co-occurrences with item1. Now, we need to "unpack" the stripe, and sort the pairs.

```
In [26]: %%writefile mapper_35b.py
#!/usr/bin/python
## mapper.py
## Author: Miki Seltzer
## Description: mapper code for HW3.5b

import sys
import itertools

# Increment mapper counter
sys.stderr.write("reporter:counter:Custom_Counter,Mapper,1\n")

# Initialize variables
total = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    # Define our key and value
    fields = line.replace('\n','').split('\t')
    word = fields[0]
    stripe = eval(fields[1])
    # Now we need to "unpack" the stripe and emit each pair for sorting
    for item in stripe:
        if stripe[item] >= 100:
            print '%s\t%s\t%s' % (stripe[item], word, item)
```

Overwriting mapper_35b.py

Since this mapper unpacks the pairs and emits them in the same format as the Pairs method in HW3.4, we can use the same reducer from the previous part to find the relative frequencies and top 50 pairs of items.

```
In [28]: # Change permissions on mapper and reducer
!chmod +x mapper_35b.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/hw3_5b_output

# Run job
!time hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFi
eldBasedComparator \
-D mapred.text.key.partitioner.options=-k1,1 \
-D stream.num.map.output.key.fields=3 \
-D mapred.text.key.comparator.options='-k1,1nr -k2,2 -k3,3' \
-D mapred.reduce.tasks=1 \
-mapper /home/cloudera/Documents/W261-Fall2016/Week03/mapper_35b.py \
-reducer /home/cloudera/Documents/W261-Fall2016/Week03/reducer_34b.py \
-input /user/miki/week03/hw3_5a_output/part* \
-output /user/miki/week03/hw3_5b_output

Deleted /user/miki/week03/hw3_5b_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/
streamjob2185279117818801498.jar tmpDir=null

real    0m34.390s
user    0m5.735s
sys     0m0.335s

In [29]: !hdfs dfs -copyToLocal /user/miki/week03/hw3_5b_output/part-00000 hw3_5b
_output.txt
```

List the top 50 product pairs


```
In [32]: print_results('hw3_5b_output.txt', 50, 0)
```

	item1	item2	count	relative freq
[1]	DAI62779	ELE17451	1,592	5.12%
[2]	FR040251	SNA80324	1,412	4.54%
[3]	DAI75645	FR040251	1,254	4.03%
[4]	FR040251	GRO85051	1,213	3.90%
[5]	DAI62779	GRO73461	1,139	3.66%
[6]	DAI75645	SNA80324	1,130	3.63%
[7]	DAI62779	FR040251	1,070	3.44%
[8]	DAI62779	SNA80324	923	2.97%
[9]	DAI62779	DAI85309	918	2.95%
[10]	ELE32164	GRO59710	911	2.93%
[11]	DAI62779	DAI75645	882	2.84%
[12]	FR040251	GRO73461	882	2.84%
[13]	DAI62779	ELE92920	877	2.82%
[14]	FR040251	FR092469	835	2.68%
[15]	DAI62779	ELE32164	832	2.68%
[16]	DAI75645	GRO73461	712	2.29%
[17]	DAI43223	ELE32164	711	2.29%
[18]	DAI62779	GRO30386	709	2.28%
[19]	ELE17451	FR040251	697	2.24%
[20]	DAI85309	ELE99737	659	2.12%
[21]	DAI62779	ELE26917	650	2.09%
[22]	GRO21487	GRO73461	631	2.03%
[23]	DAI62779	SNA45677	604	1.94%
[24]	ELE17451	SNA80324	597	1.92%
[25]	DAI62779	GRO71621	595	1.91%
[26]	DAI62779	SNA55762	593	1.91%
[27]	DAI62779	DAI83733	586	1.88%
[28]	ELE17451	GRO73461	580	1.86%
[29]	GRO73461	SNA80324	562	1.81%
[30]	DAI62779	GRO59710	561	1.80%
[31]	DAI62779	FR080039	550	1.77%
[32]	DAI75645	ELE17451	547	1.76%
[33]	DAI62779	SNA93860	537	1.73%
[34]	DAI55148	DAI62779	526	1.69%
[35]	DAI43223	GRO59710	512	1.65%
[36]	ELE17451	ELE32164	511	1.64%
[37]	DAI62779	SNA18336	506	1.63%
[38]	ELE32164	GRO73461	486	1.56%
[39]	DAI62779	FR078087	482	1.55%
[40]	DAI85309	ELE17451	482	1.55%
[41]	DAI62779	GRO94758	479	1.54%
[42]	DAI62779	GRO21487	471	1.51%
[43]	GRO85051	SNA80324	471	1.51%
[44]	ELE17451	GRO30386	468	1.50%
[45]	FR085978	SNA95666	463	1.49%
[46]	DAI62779	FR019221	462	1.49%
[47]	DAI62779	GRO46854	461	1.48%
[48]	DAI43223	DAI62779	459	1.48%
[49]	ELE92920	SNA18336	455	1.46%
[50]	DAI88079	FR040251	446	1.43%

Report the compute time for the Stripes job.

The 1st job (counts) reports the following compute times:

```
real    0m44.977s
user    0m5.111s
sys     0m0.337s
```

The 2nd job (sorts) reports the following compute times:

```
real    0m34.390s
user    0m5.735s
sys     0m0.335s
```

Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers)

Cloudera QuickStart VM: single computer, 2 processors, 2 mappers (default), 1 reducer

How many times is each mapper and reducer called?

- Mapper: 2
- Reducer: 1

Discuss the differences in these counts between the Pairs and Stripes jobs

Below is a table showing the timings for the pairs and stripes jobs (1st job counts the co-occurrences, 2nd job sorts pairs):

item	Pairs count	Stripes count	Pairs sort	Stripes sort
real	0m50.284s	0m44.977s	0m26.758s	0m34.390s
user	0m5.474s	0m5.111s	0m4.982s	0m5.735s
sys	0m0.365s	0m0.337s	0m0.345s	0m0.335s

Indeed, the stripes job took less time to complete than the pairs job in the counting phase.

However, the pairs job took less time to complete when attempting to sort the pairs. This is likely due to the fact that in order to sort the output of the stripes job, we need to "unpack" the stripes to recover each individual pair, whereas the output of the pairs job does not need any additional unpacking.


Optional problems

HW3.7. Shopping Cart Analysis

Product Recommendations: The action or practice of selling additional products or services to existing customers is called cross-selling. Giving product recommendation is one of the examples of cross-selling that are frequently used by online retailers. One simple method to give product recommendations is to recommend products that are frequently browsed together by the customers.

Suppose we want to recommend new products to the customer based on the products they have already browsed on the online website. Write a program using the A-priori algorithm to find products which are frequently browsed together. Fix the support to $s = 100$ (i.e. product sets need to occur together at least 100 times to be considered frequent) and find itemsets of size 2 and 3.

In [1]:



```

%%writefile mapper_37.py
#!/usr/bin/env python
import itertools
import sys

item_count = int(sys.argv[1])
valid_items = set()

# If our item count is greater than 1, then load the corresponding model
# file
# indicating the items we should care about.

if item_count > 1:
    model_id = str(item_count - 1)

    # The first k items in each model row will correspond to the product
    # s. We
    # can build up the set of valid items simply by iterating over the m
    # odel
    # and adding each of the elements in the first k columns.

    with open('apriori_model_' + model_id + '.txt') as model_file:
        for line in model_file:
            model_row = line.strip().split('\t')
            old_itemset = model_row[0:item_count - 1]
            valid_items.update(old_itemset)

"""
Emit all the itemsets for this basket.
"""
def emit_itemsets(basket):
    # First, we need to find out which items in the basket match the ite
    # ms
    # which we can accept in our k-itemsets. Note that we will accept ev
    # ery
    # item when the item count is 1.

    matching_items = []

    for item in basket:
        if item_count == 1 or item in valid_items:
            matching_items.append(item)

    # If we don't have enough items, we have no itemsets to emit.

    if len(matching_items) < item_count:
        return

    # Otherwise, emit all possible subsets. We'll use the pairs approach
    # to
    # make things easier to read.

    for itemset in itertools.permutations(matching_items, item_count):

```

```

        print '\t'.join(itemset) + '\t1'

    # Also emit a counter for subcombinations so that we can create a
    # tally to use for computing confidence.

    if item_count > 1:
        for sub_itemset in itertools.permutations(matching_items, item_c
ount - 1):
            print '\t'.join(sub_itemset) + '\t*\t1'

    # Finally, counter so that we can track the number of matching baske
ts.

    print ('*\t' * item_count) + str(1)


# Iterate over the baskets and emit the itemsets for each basket.

for line in sys.stdin:
    basket = line.strip().split(' ')
    emit_itemsets(basket)

```

Writing mapper_37.py

In [2]:




```

%%writefile reducer_37.py
#!/usr/bin/env python
import sys

support_threshold = int(sys.argv[1])

basket_count = 0
confidence_count = 0

current_itemset = None
current_count = 0

"""
Emit the current itemset and its count if they exceed support_threshold.
"""
def emit_count():

    # Declare that we want to use the global basket_count and confidence
    _count
    # variables rather than something local to the function.

    global basket_count, confidence_count

    # Check if we haven't started counting anything yet.

    if current_itemset is None:
        return

    # Check if we are computing the basket count from the sort operation.

    if current_itemset[0] == '*':
        basket_count = current_count
        return

    if current_itemset[-1] == '*':
        confidence_count = current_count
        return

    # Check if we have exceeded the necessary threshold.

    if current_count >= support_threshold:
        frequency = 1.0 * current_count / basket_count

        itemset_stats = str(current_count) + '\t' + str(frequency)

        if len(current_itemset) > 1:
            confidence = 1.0 * current_count / confidence_count
            itemset_stats += '\t' + str(confidence)

        print '\t'.join(current_itemset) + '\t' + itemset_stats

for line in sys.stdin:

```

```

    # Each line corresponds to the itemset stats. The last item will be
a count
    # value, while the first items will be the itemset.

    itemset_stats = line.strip().split('\t')

    itemset = itemset_stats[0:-1]
    count = int(itemset_stats[-1])

    # If we haven't switched itemsets, continue accumulating the counte
r.

    if current_itemset == itemset:
        current_count += count
        continue

    # If we have switched itemsets, emit the count for the old itemset a
nd then
    # switch to the new itemset.

    emit_count()
    current_itemset = itemset
    current_count = count

# We are guaranteed to not have printed the very last itemset, so emit i
t now.

emit count()

```

Writing reducer_37.py

Test with the sample data from the slides.

```

In [1]: !echo Beer Diaper BabyPowder Bread Umbrella > SampleSlidesData.txt
!echo Diaper BabyPowder >> SampleSlidesData.txt
!echo Beer Diaper Milk >> SampleSlidesData.txt
!echo Diaper Beer Detergent >> SampleSlidesData.txt
!echo Beer Milk CocaCola >> SampleSlidesData.txt

```

```

In [2]: !cat SampleSlidesData.txt

```

```

Beer Diaper BabyPowder Bread Umbrella
Diaper BabyPowder
Beer Diaper Milk
Diaper Beer Detergent
Beer Milk CocaCola

```

```
In [3]: !cat SampleSlidesData.txt | python mapper_37.py 1 | sort -k1 | python reducer_37.py 2 | tee apriori_model_1.txt
```

BabyPowder	2	0.4
Beer	4	0.8
Diaper	4	0.8
Milk	2	0.4

```
In [4]: !cat SampleSlidesData.txt | python mapper_37.py 2 | sort -k1 -k2 | python reducer_37.py 2 | tee apriori_model_2.txt
```

BabyPowder	Diaper	2	0.4	1.0
Beer	Diaper	3	0.6	0.75
Beer	Milk	2	0.4	0.5
Diaper	BabyPowder	2	0.4	0.5
Diaper	Beer	3	0.6	0.75
Milk	Beer	2	0.4	1.0

Test with the actual data.

```
In [5]: !cat ProductPurchaseData.txt | python mapper_37.py 1 | sort -k1 | python reducer_37.py 100 > apriori_model_1.txt
```

```
In [6]: !cat ProductPurchaseData.txt | python mapper_37.py 2 | sort -k1 -k2 | python reducer_37.py 100 > apriori_model_2.txt
```

```
In [7]: !cat ProductPurchaseData.txt | python mapper_37.py 3 | sort -k1 -k2 -k3 | python reducer_37.py 100 > apriori_model_3.txt
```

Test using Hadoop

```
In [15]: # Change permissions on mapper and reducer
!chmod +x mapper_37.py
!chmod +x reducer_37.py

# If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/aprior_model_1

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFieldBasedComparator \
-D mapred.text.key.partitionner.options=-k1 \
-D stream.num.map.output.key.fields=3 \
-D mapred.text.key.comparator.options='-k1 -k2' \
-D mapred.reduce.tasks=1 \
-mapper '/home/cloudera/Documents/W261-Fall2016/Week03/mapper_37.py 1' \
-reducer '/home/cloudera/Documents/W261-Fall2016/Week03/reducer_37.py 100' \
-input /user/miki/week03/ProductPurchaseData.txt \
-output /user/miki/week03/apriori_model_1
```

```
Deleted /user/miki/week03/hw3_7_output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/streamjob2467148671885451279.jar tmpDir=null
```

```
In [25]: # If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/apriori_model_2

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFiel
eldBasedComparator \
-D mapred.text.key.partitionner.options=-k1 \
-D stream.num.map.output.key.fields=3 \
-D mapred.text.key.comparator.options='-k1 -k2' \
-D mapred.reduce.tasks=1 \
-mapper '/home/cloudera/Documents/W261-Fall2016/Week03/mapper_37.py 2' \
-reducer '/home/cloudera/Documents/W261-Fall2016/Week03/reducer_37.py 10
0' \
-file /home/cloudera/Documents/W261-Fall2016/Week03/apriori_model_1.txt \
-input /user/miki/week03/ProductPurchaseData.txt \
-output /user/miki/week03/apriori model 2
```

```
rm: `/user/miki/week03/apriori_model_2': No such file or directory
packageJobJar: [/home/cloudera/Documents/W261-Fall2016/Week03/apriori_m
odel_1.txt] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/stream
job890159217878330682.jar tmpDir=null
```

```
In [26]: # If output folder already exists, delete it
!hdfs dfs -rm -r /user/miki/week03/apriori_model_3

# Run job
!hadoop jar hadoop-streaming-2.7.1.jar \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFiel
eldBasedComparator \
-D mapred.text.key.partitionner.options=-k1 \
-D stream.num.map.output.key.fields=3 \
-D mapred.text.key.comparator.options='-k1 -k2 -k3' \
-D mapred.reduce.tasks=1 \
-mapper '/home/cloudera/Documents/W261-Fall2016/Week03/mapper_37.py 3' \
-reducer '/home/cloudera/Documents/W261-Fall2016/Week03/reducer_37.py 10
0' \
-file /home/cloudera/Documents/W261-Fall2016/Week03/apriori_model_2.txt \
-input /user/miki/week03/ProductPurchaseData.txt \
-output /user/miki/week03/apriori_model_3
```

```
rm: `/user/miki/week03/apriori_model_3': No such file or directory
packageJobJar: [/home/cloudera/Documents/W261-Fall2016/Week03/apriori_m
odel_2.txt] [/usr/jars/hadoop-streaming-2.6.0-cdh5.5.0.jar] /tmp/stream
job3095742417588462696.jar tmpDir=null
```

```
In [28]: !hdfs dfs -cat /user/miki/week03/apriori model 3/part-00000 | head -n 20
```

DAI22896	DAI62779	GR073461	101	0.0037998495109
1	0.340067340067			
DAI22896	GR073461	DAI62779	101	0.0037998495109
1	0.332236842105			
DAI23334	DAI62779	ELE92920	143	0.0053799849510
9	0.52380952381			
DAI23334	ELE92920	DAI62779	143	0.0053799849510
9	1.0			
DAI31081	DAI62779	ELE17451	103	0.0038750940556
8	0.282967032967			
DAI31081	DAI75645	FR040251	122	0.004589917231
0.592233009709				
DAI31081	ELE17451	DAI62779	103	0.0038750940556
8	0.449781659389			
DAI31081	ELE32164	GR059710	112	0.0042136945071
5	0.358974358974			
DAI31081	FR040251	DAI75645	122	0.004589917231
0.435714285714				
DAI31081	FR040251	GR085051	102	0.0038374717833
0.364285714286				
DAI31081	FR040251	SNA80324	103	0.0038750940556
8	0.367857142857			
DAI31081	GR059710	ELE32164	112	0.0042136945071
5	0.598930481283			
DAI31081	GR085051	FR040251	102	0.0038374717833
1.0				
DAI31081	SNA80324	FR040251	103	0.0038750940556
8	0.544973544974			
DAI42083	DAI62779	DAI92600	105	0.0039503386004
5	0.905172413793			
DAI42083	DAI92600	DAI62779	105	0.0039503386004
5	0.415019762846			
DAI42083	DAI92600	ELE17451	117	0.0044018058690
7	0.462450592885			
DAI42083	ELE17451	DAI92600	117	0.0044018058690
7	0.632432432432			
DAI42493	DAI62779	ELE17451	112	0.0042136945071
5	0.363636363636			
DAI42493	DAI62779	ELE92920	112	0.0042136945071
5	0.363636363636			

cat: Unable to write to output stream.